

The document **ranking** problem

- Have a collection of documents
- User issues a query
- A list of documents needs to be returned
- **Ranking method is core of an IR system:**
 - In what order do we present documents to the user?
 - We want the "best" document to be first, second best second, etc....
- **Idea: Rank by probability of relevance of the document w.r.t. information need**
 - $P(\text{relevant} \mid \text{document}, \text{query})$

Recall a few probability basics

- For events a and b :
- Bayes' Rule

$$p(a, b) = p(a \cap b) = p(a \mid b) p(b) = p(b \mid a) p(a)$$

$$p(\bar{a} \mid b) p(b) = p(b \mid \bar{a}) p(\bar{a})$$

$$p(a \mid b) = \frac{p(b \mid a) p(a)}{p(b)} = \frac{p(b \mid a) p(a)}{\sum_{x=a, \bar{a}} p(b \mid x) p(x)}$$

← Prior

Posterior

- Odds:

$$O(a) = \frac{p(a)}{p(\bar{a})} = \frac{p(a)}{1 - p(a)}$$

Probability Ranking Principle (PRP)

Probability Ranking Principle (PRP)

Let x be a document in the collection.

Let R represent **relevance** of a document w.r.t. given (fixed) query

Let NR represent **non-relevance**.

$R=\{0,1\}$ vs. NR/R

Need to find $p(R/x)$ ---- probability that a document x is **relevant**.

$$p(R \mid x) = \frac{p(x \mid R) p(R)}{p(x)}$$

$$p(NR \mid x) = \frac{p(x \mid NR) p(NR)}{p(x)}$$

$$p(R \mid x) + p(NR \mid x) = 1$$

$p(R), p(NR)$

prior probability of retrieving a (non) relevant document

$p(x/R), p(x/NR)$

probability that if a relevant (non-relevant) document is retrieved, it is x .

Probability Ranking Principle (续)

- Simple case: no selection costs or other utility concerns that would differentially weight errors
- **Bayes' Optimal Decision Rule**
 - x is **relevant** iff $p(R|x) > p(NR|x)$
- **PRP** in action: **Rank** all documents by $p(R|x)$
- Theorem:
 - Using the PRP is optimal, in that it minimizes the loss (Bayes risk) under 1/0 loss
 - Provable if all probabilities correct, etc. [e.g., Ripley 1996]

Probability Ranking Principle (续)

- How do we compute all those probabilities?
 - Do not know exact probabilities, have to use estimates
 - **Binary Independence Retrieval (BIR)** is the simplest model
- Questionable **assumptions**
 - "Relevance" of each document is independent of relevance of other documents.
 - Really, it's bad to keep on returning **duplicates**
 - Boolean model of **relevance**
 - That one has a single step information need
 - Seeing a range of results might let user refine query

Probabilistic Retrieval Strategy

- **Estimate** how **terms** contribute to **relevance**
 - How do things like **tf**, **df**, and **length** influence your judgments about document relevance?
 - One answer is the Okapi formulae (S. Robertson)
- **Combine** to find document **relevance probability**
- **Order** documents by **decreasing probability**

Binary Independence Model (BIM)

Binary Independence Model (BIM)

- Traditionally used in conjunction with PRP
- "Binary" = Boolean:
 - documents are represented as binary incidence **vectors of terms** (cf. lecture 1):

$$\vec{x} = (x_1, \dots, x_n)$$

$x_i = 1$ iff term i is present in document x .

- "Independence": terms occur in documents independently
- Different documents can be modeled as same vector
- Bernoulli Naive Bayes model (cf. text categorization!)

Binary Independence Model (续)

- Queries: binary term incidence vectors
- Given query q ,
 - for each document d need to compute $p(R|q, d)$.
 - replace with computing $p(R|q, x)$ where x is binary term incidence vector representing d Interested only in ranking
- Will use odds and Bayes' Rule:

$$O(R|q, \vec{x}) = \frac{p(R|q, \vec{x})}{p(NR|q, \vec{x})} = \frac{\frac{p(R|q)p(\vec{x}|R,q)}{p(\vec{x}|q)}}{\frac{p(NR|q)p(\vec{x}|NR,q)}{p(\vec{x}|q)}}$$

Binary Independence Model (续)

$$O(R|q, \vec{x}) = \frac{p(R|q, \vec{x})}{p(NR|q, \vec{x})} = \frac{p(R|q)}{p(NR|q)} \cdot \frac{p(\vec{x}|R, q)}{p(\vec{x}|NR, q)}$$

Constant for a given query

Needs estimation

- Using **Independence** Assumption:

$$\frac{p(\vec{x}|R, q)}{p(\vec{x}|NR, q)} = \prod_{i=1}^n \frac{p(x_i|R, q)}{p(x_i|NR, q)}$$

• So: $O(R|q, d) = O(R|q) \cdot \prod_{i=1}^n \frac{p(x_i|R, q)}{p(x_i|NR, q)}$

Binary Independence Model (续)

$$O(R|q, d) = O(R|q) \cdot \prod_{i=1}^n \frac{p(x_i|R, q)}{p(x_i|NR, q)}$$

- Since x_i is either 0 or 1:

$$O(R|q, d) = O(R|q) \cdot \prod_{x_i=1} \frac{p(x_i=1|R, q)}{p(x_i=1|NR, q)} \cdot \prod_{x_i=0} \frac{p(x_i=0|R, q)}{p(x_i=0|NR, q)}$$

- Let $p_i = p(x_i=1|R, q)$; $r_i = p(x_i=1|NR, q)$;

- Assume, for all terms not occurring in the query ($q_i=0$) $p_i = r_i$

Binary Independence Model (续)

$$O(R|q, \vec{x}) = O(R|q) \cdot \prod_{x_i=q_i=1} \frac{p_i}{r_i} \cdot \prod_{\substack{x_i=0 \\ q_i=1}} \frac{1-p_i}{1-r_i}$$

All matching terms Non-matching query terms

$$= O(R|q) \cdot \prod_{x_i=q_i=1} \frac{p_i(1-r_i)}{r_i(1-p_i)} \cdot \prod_{q_i=1} \frac{1-p_i}{1-r_i}$$

All matching terms All query terms

Binary Independence Model (续)

$$O(R|q, \vec{x}) = O(R|q) \cdot \prod_{x_i=q_i=1} \frac{p_i(1-r_i)}{r_i(1-p_i)} \cdot \prod_{q_i=1} \frac{1-p_i}{1-r_i}$$

Constant for each query

Only quantity to be estimated for rankings

- Retrieval Status Value:

$$RSV = \log \prod_{x_i=q_i=1} \frac{p_i(1-r_i)}{r_i(1-p_i)} = \sum_{x_i=q_i=1} \log \frac{p_i(1-r_i)}{r_i(1-p_i)}$$

Binary Independence Model (续)

- All boils down to computing RSV.

$$RSV = \log \prod_{x_i=q_i=1} \frac{p_i(1-r_i)}{r_i(1-p_i)} = \sum_{x_i=q_i=1} \log \frac{p_i(1-r_i)}{r_i(1-p_i)}$$

$$RSV = \sum_{x_i=q_i=1} c_i; \quad c_i = \log \frac{p_i(1-r_i)}{r_i(1-p_i)}$$

Binary Independence Model (续)

- Estimating RSV coefficients.
- For each term i look at this table of document counts:

Documens	Relevant	Non-Relevant	Total
$X_i=1$	s	$n-s$	n
$X_i=0$	$S-s$	$N-n-S+s$	$N-n$
Total	S	$N-S$	N

- Estimates: $p_i \approx \frac{s}{S}$ $r_i \approx \frac{(n-s)}{(N-S)}$

$$c_i \approx K(N, n, S, s) = \log \frac{s/(S-s)}{(n-s)/(N-n-S+s)}$$

For now, assume no zero terms. More next lecture.

PRP and BIR

- Getting reasonable approximations of probabilities is possible.
- Requires restrictive assumptions:
 - *term independence*
 - *terms not in query don't affect the outcome*
 - *boolean representation of documents/queries/relevance*
 - *document relevance values are independent*
- Some of these assumptions can be removed
- Problem: either require partial relevance information or only can derive somewhat inferior term weights

Probabilistic Relevance Feedback

1. Guess a preliminary probabilistic description of R and use it to retrieve a first set of documents V , as above.
2. Interact with the user to refine the description :
learn some definite members of R and NR
3. Reestimate p_i and r_i on the basis of these
 - Or can combine new information with original guess (use Bayesian prior):
4. Repeat, thus generating a succession of approximations to R .

$$p_i^{(2)} = \frac{|V_i| + \kappa p_i^{(1)}}{|V| + \kappa}$$

κ is prior weight

Estimation – key challenge

If non-relevant documents are approximated by the whole collection, then

- $r_i = n_i / N$
 - r_i : prob. of occurrence in non-relevant documents for query
 - $\log(1 - r_i) / r_i = \log(N - n_i) / n_i \approx \log N / n_i = \text{IDF}$
- p_i can be estimated in various ways:
 - p_i : probability of occurrence in relevant documents
 - from relevant documents if know some
 - Relevance weighting can be used in feedback loop
 - constant (Croft and Harper combination match)
 - then just get *idf* weighting of terms
 - proportional to prob. of occurrence in collection
 - more accurately, to log of this (Greiff, SIGIR 1998)

Iteratively estimating p_i

1. Assume that p_i constant over all x_i in query
 - $p_i = 0.5$ (even odds) for any given doc
2. Determine guess of relevant document set:
 - V is fixed size set of highest ranked documents on this model
 - note: now a bit like *tf.idf*
3. Need to improve our guesses for p_i and r_i , so
 - Use distribution of x_i in docs in V . Let V_i be set of documents containing x_i
 - $p_i = |V_i| / |V|$
 - Assume if not retrieved then not relevant
 - $r_i = (n_i - |V_i|) / (N - |V|)$
4. Go to 2. until converges then return ranking

21

Probabilistic Model (summary)

- Definition
 - The index term weight : all binary , $w_{ij} \in \{0,1\}$, $w_{iq} \in \{0,1\}$
 - Query q : a subset of index terms
 - R : the set of relevant documents known (or initially guessed)
 - \bar{R} : the complement of R (the set of non-relevant documents)
 - $P(R|d_i)$: the probability that d_i is relevant to q
 - $P(\bar{R}|d_i)$: the probability that d_i is non-relevant to q
 - The similarity $sim(d_i, q)$ of d_i to q is defined as the ratio:

$$sim(d_i, q) \sim \sum_{i=1}^t w_{iq} \times w_{ij} \times \left(\log \frac{P(x_i | R)}{1 - P(x_i | R)} + \log \frac{1 - P(x_i | \bar{R})}{P(x_i | \bar{R})} \right)$$

WangWei_SEC&COSE_SEU

Assumptions

$$P(x_i | R) = 0.5, P(x_i | \bar{R}) = \frac{n_i}{N}$$

$$P(x_i | R) = \frac{V_i}{V}, P(x_i | \bar{R}) = \frac{n_i - V_i}{N - V}$$

$$P(x_i | R) = \frac{V_i + 0.5}{V + 1}, P(x_i | \bar{R}) = \frac{n_i - V_i + 0.5}{N - V + 1}$$

$$P(x_i | R) = \frac{V_i + \frac{n_i}{N}}{V + 1}, P(x_i | \bar{R}) = \frac{n_i - V_i + \frac{n_i}{N}}{N - V + 1}$$


WangWei_SEC&COSE_SEU

Bayesian Networks for Text Retrieval (Turtle and Croft 1990)

- Standard probabilistic model assumes
 - can't estimate $P(R|D, Q)$
 - Instead assume independence and use $P(D|R)$
- But maybe can with a **Bayesian network**
- What is a Bayesian network?
 - A directed acyclic graph
 - Nodes
 - Events or Variables
 - Assume values.
 - For our purposes, all Boolean
 - Links
 - model direct dependencies between nodes

Bayesian Networks

a, b, c - propositions (events).



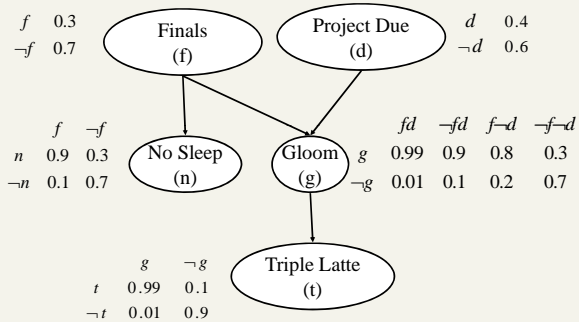
- Bayesian networks model causal relations between events

Inference in Bayesian Nets:

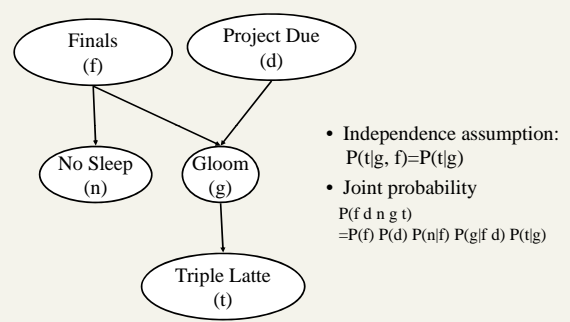
- Given probability distributions for roots and conditional probabilities can compute a priori probability of any instance
- Fixing assumptions (e.g., b was observed) will cause re-computation of probabilities

For more information see:
R.G. Cowell, A.P. Dawid, S.L. Lauritzen, and D.J. Spiegelhalter, 1999. *Probabilistic Networks and Expert Systems*. Springer Verlag.
J. Pearl. 1988. *Probabilistic Reasoning in Intelligent Systems: Networks of Plausible Inference*. Morgan-Kaufman.

Toy Example



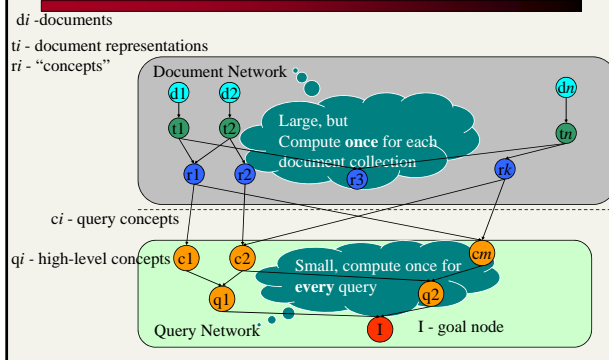
Independence Assumptions



Model for Text Retrieval

- Goal
 - given a user's information need (evidence)
 - find probability a doc satisfies need
- Retrieval model
 - Model docs in a *document network*
 - Model information need in a *query network*

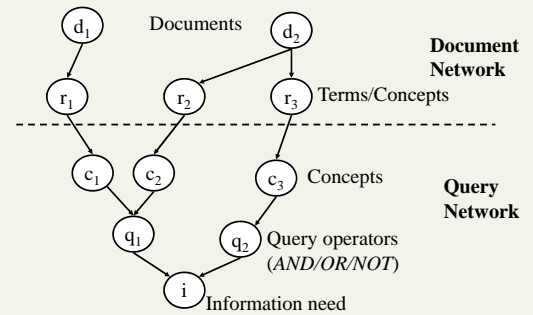
Bayesian Nets for IR: Idea



Bayesian Nets for IR

- Construct Document Network (once !)
- For each query
 - Construct best Query Network
 - Attach it to Document Network
 - Find subset of d_i 's which maximizes the probability value of node i (best subset).
 - Retrieve these d_i 's as the answer to query.

Bayesian nets for text retrieval



Link matrices and probabilities

- Prior doc probability $P(d) = 1/n$
- $P(r|d)$
 - within-document term frequency
 - $tf \times idf$ - based
- $P(c|r)$
 - 1-to-1
 - thesaurus
- $P(q|c)$
 - canonical forms of query operators
 - always use things like AND and NOT
 - never store a full CPT*

*conditional probability table

Example: "reason trouble -two"

